Predicting the Draft Slot of NBA Prospects in the Modern Era

Kevin Hillyard & Brent George
CS 472, Winter 2021
Department of Computer Science
Brigham Young University

Abstract

Every year experts attempt to determine which basketball players are the most valuable and when those players should be picked in the NBA draft. In this paper, we record our efforts to accurately predict where players would be drafted based on their amatuer statistics using different machine learning models. We first describe by what means our data was gathered and what type of data we gathered. We then review the results we had in predicting a player's draft position with our different models, namely KNN, decision tree, multi-layer perceptron, and linear regression. We then review our findings, especially from linear regression, in summarizing which statistics were important, and which were not as important. Lastly our ideas for future work and improvements to our research are made known.

1 Introduction

NBA scouts are always looking for NBA potential players. When determining how good a player is, they are almost always compared to past players. Many analysts have a general idea if a player will be a lottery pick, but the late first round and second round predictions can vary widely.

In our research we attempted to use machine learning models to predict where a player would be drafted. These models use the past 11 years of NBA draft results, along with the players' associated statistics, to determine where the player should be drafted, if they are to be drafted at all.

The models we used were Decision Tree, KNN-Regressor, Linear Regression, and Multi-Layer Perceptron (MLP). For each of these models we trained them on three different datasets with a variety of hyperparameters. The three datasets we used were a variation of our entire dataset. The first dataset only included players that were actually drafted, the second included all players, both drafted and undrafted, while the third included both drafted and undrafted players but without features that were deemed less important.

There was some difficulty in predicting a player's draft position as it is a human perception of that player's capability and is often swayed by what each NBA teams' current needs are. A possession-based point guard may have the background of a #15 pick, but if they meet a need of a team picking #5, they may be drafted at that slot. Additionally, human general managers can make baffling decisions. In 2012, Anthony Bennett was regarded as a late first round pick by the majority of the NBA, perhaps a fringe lottery pick. However, the Cleveland Cavaliers, who owned the #1 pick in the draft, selected him at that slot. He proved to be one of the biggest busts in NBA history, failing to even make an NBA roster only three years after being the first draft pick. Any machine learning model would have difficulty predicting his true draft slot.

The size of our dataset is inherently limited by the size of the NBA draft. We were also limited to a certain number of drafts as the era of basketball has changed, encouraging more long range field goals compared to years prior. This drastically affects the typical stats for players making comparisons between newer and older players less viable.

Despite those difficulties we were able to predict a player's draft position to within 10 picks. In a future study if we were to account for the aforementioned issues the accuracy of our results could be increased.

2 Methods

2.1 Data Source

There are various online providers of basketball statistics. Official sources include NBA.com and the official site of NCAA basketball. However, official sources often do not include more than the most basic statistics.

Several sites exist that provide extensive college statistics. Many of these have a paywall (kenpom.com) or do not include statistics for international prospects. Other sites request not to be scraped by bots.

After further exploration, we found realgm.com, a site that consented to limited scraping that also included historical data for past statistics, data on inter-

national prospects, and a wide array of advanced statistics. We wrote a web scraper in Python to traverse the element tree of the web page and lift each of the players selected in the past 11 NBA drafts and every available statistics associated with those 660 players. This formed our initial dataset, which was stored in a CSV file. We also pulled the records of nearly 736 prospects that went undrafted during that same time period.

2.2 Featurization of Data

Due to the nature of basketball statistics, the data scraped from the website was already largely in the form we needed. We wrote a Python script to perform a few sanitizing steps and convert the CSV file to an ARFF file.

The main sanitization we needed to perform was the consolidation of various leagues that prospects played in. Prospects played in over 35 leagues and countries. While there was consistency for prospects that played in the NCAA in the United States, there was extreme fragmentation among the teams and leagues that international prospects played for. In order to help the model, we grouped all international teams under one label, leaving us with three total league labels: NCAA, G-League, and International.

While drafted players were labeled according to their draft selection, the undrafted players that we pulled were unlabeled. We didn't want to use a clustering algorithm to label these instances - a clustering algorithm would attempt to label these players with a pick number from 1 to 60, whereas we wanted to be able to predict if a player went undrafted. We decided to label any undrafted player as being picked 61st, with the plan of interpreting our future model's output under this assumption. This led to over half of our dataset consisting of players "selecting 61st in the draft". This would prove to have some positive effects, such as giving the model many examples of poor prospects on the higher end of the draft. It would also have negative effects, such as leading the models to be biased towards higher pick numbers when it was unsure.

Most of our chosen features were already real values, and directly translated to our final dataset. We had two nominal features, position and league. We decided to do a binary encoding for each of these two features, trusting that higher complexity models would be able to interpret this and differentiate the value of various statistics per position. For example, we would anticipate that the model would value rebounds and blocks for a center more than for a point guard. This may not be a necessary distinction in the end, but we wanted to give our models the chance to pick up on it.

2.3 Models

We used a variety of models in this project, to be described in further detail in results sections. All implementations came from scikit learn v1.0.1.

3 Initial Results

Initially, we trained Decision Tree, Multi-Layer Perceptron, KNN-Neighbor, and Linear Regression models. Each of these had a grid of hyperparameters that was exhaustively searched to find the best results, as recorded in Tables 1, 2, and 3.

3.1 Without Undrafted Players

We first trained our models on the dataset excluding all undrafted players. The Mean Absolute Error (MAE) was fairly high, with an average of about 15 across all hyperparameter combinations, and with minimums of about 12 picks off for the best model and hyperparameter combinations. An issue we found with this subset was that there were not enough bad players for the models to learn form.

3.2 Decision Tree

After training the decision tree model on this subset it achieved a minimum MAE of 13.689. The hyperparameters with which we achieved the best score used poisson as the criterion, a max depth of 8 layers, minimum samples set to 6 and 4 for split and leaf respectively, and both the minimum weight fraction leaf and minimum impurity decrease set to 0.01.

3.3 KNN-Regressor

When training our KNN-Regressor on the subset excluding undrafted players it achieved a minimum MAE of 13.129. We achieved the best MAE when setting k to 9 and using a manhattan distance metric.

3.4 Linear Regression

The MAE our Linear Regression model achieved was a value of 11.265. No special hyperparameters were used for this model.

3.5 MLP

The last model which we trained on this subset of data was our MLP. It achieved a MAE of 11.490. The combination of hyperparameters that achieved the best MAE had hidden nodes set to (64, 64), a learning rate of 0.2, a momentum of 0.5, and regularization of 0.0001.

3.6 With Drafted and Undrafted Players

As already mentioned the following models achieved a better MAE when using both drafted and undrafted players. The variety of talent allowed the models to learn to recognize both good and bad players and was more accurate as a result.

3.7 Decision Tree

When training the decision tree on both drafted and undrafted players it achieved a MAE of 9.560. This is an improvement from 13.689 MAE that was achieved using the subset without undrafted players. This score

Hyperparameter	Opt. 1	Opt. 2	Opt. 3	Opt. 4	Opt. 5
Criterion	Squared Error	Friedman MSE	Absolute Error	Poisson	
Splitter	Best				
Max Depth	8	16	24	32	48
Min Samples for Split	2	4	6		
Min Samples per Leaf	1	2	4		
Min Weight Fraction per Leaf	0.0	0.001	0.01		
Min Impurity Decimal	0.0	0.001	0.01		

Table 1: Decision Tree Hyperparameter Grid

Hyperparameter	Opt. 1	Opt. 2	Opt. 3	Opt. 4	Opt. 5
Hidden Nodes	(64)	(128)	(64, 64)	(64, 32, 64)	(128, 128)
Learning Rates	0.01	0.1	0.2	0.5	1.0
Momentum	0.0	0.3	0.5	0.8	1.0
Regularization	0.0001	0.0005	0.001		

Table 2: Multi-Layer Perceptron Hyperparameter Grid

would prove to be our best overall score in our initial training. We used the same hyperparameters mentioned before and the best score was achieved when using the following settings: absolute error as the criterion, max depth set to 8 layers, the minimum samples for splits set to 4 while the minimum samples for leaves was set to 1, the minimum weight fraction leaf set to 0.01, and the minimum impurity decrease set to 0.001.

3.8 KNN-Regressor

The KNN-Regressor when trained on both drafted and undrafted players achieved a MAE of 11.699. This is 1.43 picks better than the previous KNN-Regressor model. This score would be the worst score for all models trained on both drafted and undrafted players. The score was achieved using the same hyperparameters as before but with a k value of 3 using a euclidean distance metric.

3.9 Linear Regression

Our Linear Regression model had the third best initial score with a MAE of 10.091. This is an improvement from the previous Linear Regression model that scored 11.265.

3.10 MLP

Lastly our MLP had the second best initial score with a MAE of 9.851. It used the same hyperparameters as the previous MLP model, but had the following settings: hidden nodes was set to 128, learning rate was set to 0.1, momentum was set to 0.8, and we used a regularization value of 0.0001. Compared to the first MLP model not trained on undrafted players, this model improved that score by 1.588.

3.11 Summary of Initial Results

Figure 1 shows a visual for the best results for our initial models. The models trained with the undrafted players included performed much better on average. However, there wasn't much differentiation between the MLP, Decision Tree, and Linear Regression models with that dataset, suggesting that none had a particular advantage for our compiled datasets.

4 Feature and Model Improvements

To improve our results we reviewed the most important features and least important features as determined by the weights from our linear regression model. The features that had the lowest weights included the following features: Usage %, Assist %, Field Goals Made, Field Goals Attempted, Free Throws Made, Free Throws Attempted, 3-pointers Made, and 3-pointers Attempted. We deemed these features safe to remove as they had low weights and the essence of their value was captured in a different feature. For example, Field Goals

Hyperparameter	Opt. 1	Opt. 2	Opt. 3	Opt. 4
K	1	3	5	7
K continued	9	11	13	15
Distance Metrics	Euclidean	Manhattan	Chebyshev	Minkowski

Table 3: K-Nearest Neighbors Hyperparameter Grid

Figure 1: Bar Chart of Mean Absolute Error for Each Initial Model

Figure 2: Bar Chart of Mean Absolute Error Including Smaller Dataset Results

Made and Field Goals Attempted is captured in the feature Field Goal % which had a high weight. In removing these somewhat redundant features we hoped to slightly improve our predictions, especially for the KNN model. We didn't weight the values of various features when calculating the distance between neighbors, so if we had many redundant or useless features, they would drown out the important features and make it difficult for the KNN model to be effective.

5 Final Results

5.1 Decision Tree

When training a Decision Tree model on this new dataset we were able to achieve a MAE of 9.718. This was the best score from all of the models trained on this new dataset but was 0.158 higher than the Decision Tree trained on the unaltered dataset with both drafted and undrafted players. This would be the second best score from all of our results. This score was achieved using the following hyperparameters: a maximum depth of 8 layers, a minimum sample of 2 for splits and 1 for leaves, minimum weight fraction leaf of 0.01, and a minimum impurity decrease of 0.01.

5.2 KNN-Regressor

The KNN-Regressor trained on the new dataset had the worst score of all models trained on this dataset. It had a MAE of 11.790. This was better than the first model trained on the data without undrafted players, but worse than the model trained on all players. This score was achieved using a k value of 2 and using a euclidean distance metric.

5.3 Linear Regression

Training the Linear Regression model on this new dataset resulted in a MAE of 9.987. This was the second worst of all models trained on this dataset but was better than the other two Linear Regression models trained on the other datasets.

5.4 MLP

Lastly our MLP trained on this new dataset resulted in a MAE of 9.858. This was only 0.007 worse than the model trained on the unaltered dataset with all players. This was the second best score of all models trained on the new dataset and the fourth best overall. The score was achieved by using the following hyperparameters: the hidden nodes were set to (32,32), the learning rate was 0.1, momentum was set to 0.8, and we used a regularization value of 0.001.

5.5 Improvement Results

The changes we made with our improvements initially surprised us as the MAE for each model trained on the adjusted dataset was almost identical to, if not slightly higher than, the MAE scores for the models trained on the unaltered dataset with both drafted and undrafted players. Reviewing these results they do match up with the behavior of the models we chose. All of our models, aside from the KNN-Regressor, are capable of learning to ignore unnecessary features. By removing features the linear regressor deemed less useful we most likely only reduced the training time for our models. Concerning our KNN-Regressor the consistently high MAE was likely due to not having enough data to effectively make predictions from neighboring instances. This will be discussed further in the Final Results section.

6   Discussion and Conclusion

7   Future Work

There are some minor changes we would make to our research to improve the results of our predictions. These changes include further research about the data, how the error of our models was determined, and some additional features.

7.1   Further Draft Research

As previously mentioned in this report there were some contributing factors to a player's draft stock that we did not take into account when gathering our data. When gathering our data and when a player was drafted it would be worthwhile to research the team's needs that drafted that player. Knowing this would allow us to determine if the player was drafted at that position for their skills as a player or more for their fit with that team.

Additionally, we could do player comparisons with existing NBA players. This is a common draft tactic used by general managers, and it might be helpful for the models as well. This would involve gathering player profiles of successful and unsuccessful NBA players, and using their similarity to NBA prospects to determine if that prospect might be valued in the modern NBA.

Finally, we could improve the size of our dataset by including more drafts. This is initially unhelpful because the shifting nature of the NBA leads to certain players being very valuable 40 years ago and insignificant today, and vice versa. We would have to do era adjustments on any datasets that included older generations and thoroughly test the effectiveness of this approach.

7.2   Error Scoring

There was an issue with the error scoring with our models as we were using the scikit models with their own error scoring function. Because undrafted players don't have a draft position associated with them we imputed their labels so that all undrafted players were drafted 61st.

The issue appears when our models predict a draft position greater than 61. Our data has a max of 61 while our models are able to predict draft positions higher than this. The error scoring for the models considers this difference as an error, when in reality if the player has a draft position of 61 any placement greater than or equal to 61 should be considered correct. A custom error scoring module that plugged in to skicit learns models could dramatically improve our results.

7.3   Additional Features/Sources

As a final change to our research we would gather data from one or more additional sites in order to gather all advanced statistics possible. The site we pulled our data from had a large number of advanced statistics but not all of them.